

Final Project Report: Smart Customer Support System

Natural Language Processing Laboratory (Assignment 5)

1. Abstract

In the modern digital service landscape, customer support teams are often overwhelmed by the sheer volume of incoming queries. This project presents an **End-to-End Natural Language Processing (NLP) Pipeline** designed to automate the initial triage of support tickets. By implementing a dual-model architecture, the system simultaneously categorizes the *topic* of a ticket (Billing vs. Technical) and analyzes the *sentiment* (Positive vs. Negative). The result is an intelligent routing system that can flag urgent issues from angry customers while automatically organizing standard queries, thereby reducing response times and human error.

2. Problem Statement & Motivation

The Challenge

Manual ticket sorting is inefficient. A human agent spends valuable time reading an email just to decide which department should handle it. Furthermore, amidst thousands of "How do I reset my password?" emails, a critical complaint from a furious customer might go unnoticed for hours, leading to churn.

The Solution

We propose an automated "Smart Triage" system. The goal is not to replace human agents, but to empower them by:

1. **Filtering:** Instantly routing tickets to the correct department (Finance, IT, or General).
2. **Prioritizing:** acting as a "traffic light" system where angry customers get a "Red Light" (Escalation) and standard queries get a "Green Light" (Auto-reply).

3. System Design & Methodology

3.1 Dataset Generation

Due to privacy concerns with real-world support emails, a **synthetic dataset** was programmatically generated for this project. The dataset consists of 300 labeled examples across three categories:

- **Billing:** Keywords like *invoice, charge, refund, credit card*.
- **Technical Support:** Keywords like *crash, error, login, install, server*.
- **General Inquiry:** Keywords like *hours, location, partnership, policy*.

3.2 Data Preprocessing

Raw text is noisy and unstructured. To prepare the data for modeling, a preprocessing pipeline was implemented using the NLTK library:

- **Noise Removal:** Regular expressions (re) were used to strip punctuation and special characters, leaving only alphabetic text.
- **Normalization:** All text was converted to lowercase to ensure consistency (e.g., treating "Invoice" and "invoice" as identical).
- **Stop-word Removal:** Common English words (e.g., "the", "is", "and") were removed to reduce feature space dimensionality.
- **Lemmatization:** Words were reduced to their root forms (e.g., "crashing" \rightarrow "crash") using the WordNet Lemmatizer.

3.3 Feature Engineering (TF-IDF)

Instead of simple word counts (Bag of Words), we utilized **Term Frequency-Inverse Document Frequency (TF-IDF)**. This method was chosen because support tickets often contain generic polite words ("please", "help"). TF-IDF down-weights these common terms and highlights domain-specific keywords like "invoice" or "bug," which are crucial for classification.

3.4 Model Architecture

The core innovation of this project is the use of **two parallel Logistic Regression classifiers**:

1. **The Topic Model:** Predicts the Department (Billing / Tech / General).
2. **The Sentiment Model:** Predicts the Urgency (Positive / Neutral / Negative).

Logistic Regression was selected over deeper neural networks because, for text classification with limited data, it offers an excellent balance of speed, interpretability, and performance.

4. Implementation Details

The system functions as a live command-line interface (CLI). When a user inputs a ticket, the system processes it in real-time using the following logic flow:

1. **Input:** User types "I want my money back!"
2. **Vectorization:** The text is cleaned and converted to a numeric vector.
3. **Inference:** Both models make a prediction.
4. **Routing Logic (The "Traffic Light" System):**
 - **Red Flag:** If Sentiment is *Negative* Escalate immediately (overrides topic).
 - **Yellow Flag:** If Topic is *Billing* Route to Finance.
 - **Blue Flag:** If Topic is *Tech Support* Route to IT.
 - **Green Flag:** If Topic is *General* Send automated FAQ.

5. Results & Evaluation

5.1 Performance Metrics

The models were evaluated on a 20% test split. The system achieved high accuracy (>90%) across both tasks.

- **Topic Classification:** The model showed near-perfect separation between categories. This is expected, as the vocabulary difference between "money/invoice" and "crash/login" is distinct.
- **Sentiment Analysis:** Performance was strong for clearly negative words ("hate", "terrible"), though the model showed slightly lower recall for "passive-aggressive" complaints that lacked strong negative keywords.

5.2 Visualization

The confusion matrix reveals that the primary source of error was between "General Inquiries" and "Tech Support" when the user used vague language (e.g., "How does this work?").

6. Challenges & Limitations

While the system performs well in a controlled environment, several challenges were identified during development:

1. **The Sarcasm Blind Spot:** The current TF-IDF approach struggles with sarcasm. A review stating "*Great job deleting my data!*" is classified as **Positive** because of the words "Great job," whereas a human would instantly recognize it as **Negative**.
2. **Context Limitations:** The model treats sentences as a "bag of words." It does not understand sequence. "Not good, but bad" and "Not bad, but good" look very similar to the model mathematically.
3. **Data Imbalance:** In a real-world scenario, 80% of tickets might be simple password resets. The model might become biased toward the majority class without careful resampling techniques.

7. Future Scope

To evolve this from a mini-project to a production-ready system, the following improvements are proposed:

- **Deep Learning (BERT):** Replacing Logistic Regression with a Transformer model (like BERT) would solve the sarcasm and context issues, as BERT understands the *relationship* between words.
- **Named Entity Recognition (NER):** Integrating an NER module to automatically extract Order IDs, Dates, and Email Addresses from the ticket body.
- **Web Deployment:** Wrapping the Python script in a **Streamlit** or **Flask** web interface to allow non-technical support staff to use the tool via a browser.

8. Conclusion

This project successfully demonstrated the power of NLP in automating business workflows. By combining text preprocessing, TF-IDF vectorization, and logical routing, we built a system that not only saves time but also prioritizes customer satisfaction. It serves as a strong

foundation for understanding how AI can be applied to solve real-world communication problems.